

Linux Systems Internals and Programming (contd...)

2. Segmentation and Paging

2.1 Role of Segmentation

Definition:

Segmentation is a memory management technique that divides a program's memory into variable-sized segments. Each segment corresponds to a logical part of the program, such as code, data, stack, or heap.

Key Features:

- Provides **logical division** of memory.
- Each segment has a **base address** and **limit** (size).
- Segments allow **protection** by setting different access permissions.

Real-world Example:

Consider a C program with the following memory segments:

- **Code Segment (Text Segment)** – Stores the compiled machine code.
- **Data Segment** – Contains global and static variables.
- **Stack Segment** – Stores local variables and function call information.
- **Heap Segment** – Used for dynamic memory allocation.

Why is segmentation useful?

Segmentation enables efficient memory allocation and protection by restricting access to certain regions (e.g., user processes cannot modify the kernel segment).

2.2 Paging Mechanisms for Memory Isolation

Definition:

Paging is a memory management scheme that eliminates fragmentation by dividing memory into fixed-sized blocks called **pages** (in RAM) and **frames** (in physical memory).

Mechanism:

1. The process address space is divided into **pages** of equal size (e.g., 4 KB).
2. The physical memory is divided into **frames** of the same size.
3. The OS maintains a **Page Table** that maps virtual page numbers (VPNs) to physical frame numbers (PFNs).

Key Benefits:

- ✓ No external fragmentation (unlike segmentation).

- ✓ Enables **virtual memory**, allowing processes to run even if they are not fully in RAM.
- ✓ Provides **memory isolation**, preventing one process from accessing another's memory.

Practical Example – Page Tables in Linux

Linux maintains **multi-level page tables**:

- **Page Global Directory (PGD)** → Top-level table.
 - **Page Middle Directory (PMD)** → Intermediate table.
 - **Page Table Entry (PTE)** → Maps a virtual page to a physical frame.
-

2.3 Address Translation by the MMU

Memory Management Unit (MMU):

The MMU is a hardware component that performs **address translation** from virtual addresses (used by applications) to physical addresses (used by hardware).

Steps in Address Translation:

1. CPU generates a **virtual address** (VA).
2. MMU looks up the **Page Table** to find the corresponding **physical address** (PA).
3. If the page is **not present in memory** (Page Fault), the OS loads it from disk.
4. The MMU updates the **TLB (Translation Lookaside Buffer)** for fast lookups.

Practical Example in C – Simulating Paging:

```
#include <stdio.h>

#define PAGE_SIZE 4096 // 4 KB pages
#define FRAME_COUNT 8 // Physical memory has 8 frames

int page_table[FRAME_COUNT] = {2, 4, 7, 1, 5, 3, 6, 0}; // Mapping of virtual pages to
// physical frames

void translate_address(int virtual_page) {
    if (virtual_page >= FRAME_COUNT) {
        printf("Page Fault! Virtual page %d not found.\n", virtual_page);
        return;
    }
    int physical_frame = page_table[virtual_page];
    printf("Virtual Page %d → Physical Frame %d\n", virtual_page, physical_frame);
}

int main() {
    translate_address(3);
    translate_address(5);
    translate_address(10); // Simulating a page fault
    return 0;
}
```

}

📌 Output Example:

Virtual Page 3 → Physical Frame 1
Virtual Page 5 → Physical Frame 3
Page Fault! Virtual page 10 not found.

3. User-Space to Kernel-Space Transition

3.1 Transition Mechanism

3.1.1 Understanding User-Space to Kernel-Space Transition

Definition:

A system transitions from **user mode** (restricted operations) to **kernel mode** (privileged operations) when an application needs access to hardware or sensitive resources.

How does this happen?

- **System Calls (Syscalls)** – The primary method for safe transition.
 - **Interrupts & Exceptions** – Hardware triggers requiring OS intervention.
 - **Traps & Faults** – Special CPU instructions that cause controlled transitions.
- ◆ **Example:** When a program executes `printf()`, it internally calls the **write syscall**, which interacts with the OS to write data to the terminal.
-

3.1.2 Role of System Calls (Syscalls)

Definition:

A system call is a controlled way for user applications to request privileged operations from the OS.

Steps in a Syscall Execution:

1. The application invokes a library function (e.g., `open()` in `stdio.h`).
2. The function triggers a **software interrupt** (e.g., `int 0x80` in x86, `syscall` in x86_64).
3. The CPU switches to **kernel mode**, and the OS handles the request.
4. The OS performs the operation (e.g., reading a file).
5. Control returns to the user-space application.

Practical Example – Calling a Syscall in C:

```
#include <stdio.h>
#include <unistd.h>

int main() {
    char buffer[50];
    int bytes = read(STDIN_FILENO, buffer, sizeof(buffer));
    write(STDOUT_FILENO, buffer, bytes);
    return 0;
}
```

Here, `read()` and `write()` use **syscalls** internally.

3.2 Privilege Escalation

3.2.1 How Syscalls Enable Privilege Escalation

Privilege Escalation:

This occurs when a process gains higher privileges than intended (e.g., a user process obtaining root privileges).

How it happens?

1. **Legitimate Escalation** – Running `sudo` to gain admin access.
2. **Exploit-Based Escalation** – Using buffer overflow attacks or misconfigured permissions.

♦ Example:

```
sudo su # Elevates to root privileges
```

3.2.2 Syscall Handling in the Kernel

How Linux Handles a Syscall?

1. The syscall number is stored in **registers** (e.g., `rax` in x86_64).
2. The CPU executes the **syscall instruction**, transferring control to the kernel.
3. The **system call dispatcher** looks up the syscall table.
4. The corresponding kernel function executes the requested operation.
5. The result is returned to the user process.

Example – Listing Syscalls in Linux:

```
cat /proc/syscall
strace ls # Shows syscalls used by 'ls' command
```

Additional Resources

📌 **Reference Books & Articles:**

1. "**Understanding the Linux Kernel**" – Daniel P. Bovet, Marco Cesati
2. "**Operating System Concepts**" – Abraham Silberschatz

📌 **Online Resources:**

- Linux Kernel Docs: <https://www.kernel.org/doc/>
 - OSDev Wiki: <https://wiki.osdev.org>
 - IBM Developer: <https://developer.ibm.com>
-